Suhas Kotha

kotha@stanford.edu | homepage

Education

Stanford University, PhD Computer Science — 2024 - current

Advised by Percy Liang and Tatsunori Hashimoto

Carnegie Mellon University, MS (funded) + BS Computer Science — 2020 - 2024

Advised by Aditi Raghunathan

Work experience

Thinking Machines, Research Intern (hosted by Kevin Lu and John Schulman) — Summer 2026

Jane Street, Quantitative Trading Intern — Summer 2022

NASA Frontier Development Lab, Machine Learning Intern — Summer 2020

Teaching

- Intructor/co-creator for Deep Learning Alchemy — Fall 2026
- Instructor for Hype for Types — Spring 2022, Fall 2022
- TA for Functional Programming — Spring 2021, Spring 2022
- TA for Algorithm Design/Analysis — Fall 2022
- TA for Mathematical Foundations for CS — Fall 2021

Publications

Data-efficient pre-training by scaling synthetic megadocs
Konwoo Kim*, **Suhas Kotha***, Yejin Choi, Tatsunori Hashimoto, Nick Haber, Percy Liang
COLM 2026

Replaying pre-training data improves fine-tuning
Suhas Kotha, Percy Liang
Preprint

Synthetic Mixed Training: Scaling Parametric Knowledge Acquisition Beyond RAG
Seungju Han, Konwoo Kim, Chanwoo Park, Benjamin Newman, **Suhas Kotha**, Jaehun Jung, James Zou, Yejin Choi
Preprint

Pre-training under infinite compute
Konwoo Kim*, **Suhas Kotha***, Percy Liang, Tatsunori Hashimoto
ICLR 2026, ***Oral***

Repetition Improves Language Model Embeddings
Jacob Springer, **Suhas Kotha**, Daniel Fried, Graham Neubig, Aditi Raghunathan
ICLR 2025

Jailbreaking is Best Solved by Definition

Taeyoun Kim*, **Suhas Kotha***, Aditi Raghunathan
Neurips SafeAI Workshop 2024

A Safe Harbor for AI Evaluation and Red Teaming
Shayne Longpre et al (23 authors)
*ICML 2024, **Oral** Position Paper*

Understanding Catastrophic Forgetting in Language Models via Implicit Inference
Suhas Kotha, Jacob Springer, Aditi Raghunathan
ICLR 2024

Provably Bounding Neural Network Preimages
Suhas Kotha*, Christopher Brix*, Zico Kolter, Krishnamurthy Dvijotham†, Huan Zhang†
*NeurIPS 2023, **Spotlight***

Talks

- ➢ Data-efficient pre-training by scaling synthetic megadocs
  - DatologyAI
  - Apple Intelligence
  - Cohere Research
  - CS525 Training Data for AI
- ➢ Pre-training under infinite compute
  - Carnegie Mellon University
  - ICLR 2026 Oral
  - Cursor Research
  - Jane Street ML Research
  - Stanford NLP 25 Year Reuniuon
  - ASAP Seminar
  - Noah Goodman Group
  - Brian Trippe Group
  - Ludwig Schmidt Group
  - Surya Ganguli Group
  - Diyi Yang Group
  - Stanford NLP Lunch
- ➢ Replaying pre-training data improves fine-tuning
  - DatologyAI
  - Test-Time Training Research Institute
- ➢ Understanding Catastrophic Forgetting in Language Models via Implicit Inference
  - Brave Machine Learning Research
- ➢ Provably Bounding Neural Network Preimages
  - UIUC Formal Verification Group
  - CMU Undergraduate Theory Community